

SQIsign v2.0.1 Specification Review

Feedback from an Independent Rust Implementation

Deirdre Connolly

Selkie Cryptography

durumcrustulum@gmail.com

April 16, 2026 (living document)

Abstract

Each item below describes something that the SQIsign specification [1] (version 2.0.1, dated 2025-07-07) leaves unstated, ambiguous, or implicit, and that we could only resolve by inspecting the C reference implementation. We file these as constructive suggestions so that future implementors can achieve interoperability from the spec alone.

Contents

1	Severity scale	2
2	$\mathbb{F}_{p^2}$ square root canonicalization	2
3	ProductToTheta coordinate convention	2
4	“Squared theta” definition	3
5	hadamard_bool in the (2, 2)-chain	3
6	Torsion basis slot convention	3
7	Normalization of $(A_{24} : C_{24})$	3
8	Gluing requires Jacobian y-coordinates	4
9	Jacobian doubling for balanced strategy	4
10	DIFFERENCEPOINT normalization factor	4
A	Never recompute $P-Q$ via DIFFERENCEPOINT	5
B	Change-of-basis matrix application convention	5
C	GENERALIZEDREPRESENTINTEGER sampling ranges and divisibility	6
D	Response lattice: product vs. intersection	7
E	FIXEDDEGREEISOGENY missing u^{-1} normalization	7
F	Fixed-precision Hermite Normal Form	7

G	Endomorphism image $P-Q$ and the Okeya-Sakurai lift	8
H	GETINDEXSPLITTING must handle malformed input	9
I	IDEALTOISOGENY line 6: matrix formula is under-specified	9
J	IDEALTOISOGENY: β lives in the reduced ideal, not the input I	10
K	SUITABLEIDEALS line 14: $v_2(u)$ vs. $v_2(\gcd(u, v))$	11
L	Algorithm 8.47 implementation variants are not specified	11

1 Severity scale

CRITICAL	The spec as written produces a non-interoperable implementation: verification rejects valid signatures, or signing produces unverifiable ones. The spec must be amended.
HIGH	An implementor will almost certainly write a bug that passes unit tests but fails KAT or cross-implementation testing. Debugging requires reading the C reference source.
MEDIUM	An implementor may write a bug depending on reasonable but wrong assumptions about conventions. A short note eliminates the ambiguity.
Low	The spec is correct but incomplete. A remark or forward pointer would help.

2 $\mathbb{F}_{p^2}$ square root canonicalization

Affects Algorithm 8.12 (DIFFERENCEPOINT), and transitively every algorithm that consumes a torsion basis. **CRITICAL**

Algorithm 8.12 says “compute $\delta = \sqrt{B_{XZ}^2 - B_{XX}B_{ZZ}}$ ” without specifying which of the two square roots in $\mathbb{F}_{p^2}$ to use. The two roots δ and $-\delta$ produce two *affinely distinct* points x_{P-Q} . These are not projectively equivalent; they are genuinely different points. The wrong choice cascades through LADDER3PT (wrong kernel generator), the challenge isogeny (different curve E_{chl}), and every subsequent verification step. In our implementation the j -invariant of E_{chl} was completely wrong until we matched the C reference’s canonicalization.

The spec should define a canonical $\mathbb{F}_{p^2}$ square root. The C reference uses the convention from Aardal et al. [2]: negate the output so that the real part (as an integer in $[0, p)$) is even; if the real part is zero, negate so the imaginary part is even. This should be stated in §8.1 (field arithmetic) as a normative requirement. A note of the form

Throughout this document, “compute $\sqrt{a}$ ” for $a \in \mathbb{F}_{p^2}$ means the unique root r such that $r^2 = a$ and $\text{Re}(r) \bmod 2 = 0$ (or, if $\text{Re}(r) = 0$, $\text{Im}(r) \bmod 2 = 0$).

would suffice.

3 ProductToTheta coordinate convention

Affects Algorithm 8.37 (PRODUCTTOTHETA).

HIGH

The product theta coordinates are described abstractly via level-2 theta functions. The mapping from Montgomery projective coordinates $(X : Z)$ to product theta coordinates is not stated explicitly. The natural “theta-like” representation for Montgomery arithmetic is $(X - Z, X + Z)$, which appears

throughout the isogeny literature. The C reference uses $\text{raw}(X, Z)$ directly. Using $(X - Z, X + Z)$ produces a completely wrong change-of-basis matrix N .

Algorithm 8.37 should state explicitly that the product theta coordinates are formed as

$$(a, b, c, d) = (X_1X_2, X_1Z_2, Z_1X_2, Z_1Z_2)$$

from Montgomery projective pairs $(X_1 : Z_1), (X_2 : Z_2)$.

4 “Squared theta” definition

Affects Algorithm 8.29, Algorithm 8.34, Algorithm 8.38.

HIGH

Several algorithms refer to “squared theta coordinates” without a centralized definition. In the C reference, `to_squared_theta(P)` computes component-wise squaring *followed by a Hadamard transform*: $\tilde{P} = \mathcal{H}(P^2)$. The name suggests only squaring. We omitted the inner Hadamard in the generic isogeny evaluation, computing $\mathcal{H}(\text{inv} \cdot P^2)$ instead of the correct $\mathcal{H}(\text{inv} \cdot \mathcal{H}(P^2))$.

Add a definition early in §8.5:

*The **squared theta** of a point $P = (a : b : c : d)$ is $\tilde{P} = \mathcal{H}(a^2, b^2, c^2, d^2)$, i.e. component-wise squaring followed by the Hadamard transform.*

5 hadamard_bool in the (2, 2)-chain

Affects Algorithm 8.47 (ISOGENY22CHAINWITHTORSION).

HIGH

The algorithm describes each generic step uniformly. The C reference uses three formula variants controlled by boolean flags:

- Normal ($i < n-2$): codomain Hadamard-transformed; $\text{eval} = \mathcal{H}(\text{precomp} \cdot \mathcal{H}(P^2))$.
- Penultimate ($i = n-2$): codomain in dual form; $\text{eval} = \text{precomp} \cdot \mathcal{H}(P^2)$.
- Ultimate ($i = n-1$): extra input Hadamard, codomain dual; $\text{eval} = \text{precomp} \cdot \mathcal{H}(\mathcal{H}(P)^2)$.

The splitting step expects dual form. Using the normal formula for all steps feeds the splitter a Hadamard-transformed null point with no zero $U_{i,j}(0)$ entry, causing verification to fail after an otherwise correct 125-step chain.

A short remark in Algorithm 8.47 noting that the last two steps omit the final Hadamard (and the ultimate step pre-Hadamards the input) would prevent this.

6 Torsion basis slot convention

Affects Algorithm 2.2 (TORSIONBASISFROMHINT).

MEDIUM

The spec says the algorithm returns $(P, Q, P-Q)$. The C reference stores $P-Q$ in the Q slot and Q in PmQ , so LADDER3PT computes $P + [m](P-Q)$, not $P + [m]Q$. The swap ensures the multiplied point avoids $(0, 0)$ but is not documented.

State the output convention explicitly, or note that LADDER3PT is called with the difference in the Q position.

7 Normalization of $(A_{24} : C_{24})$

Affects Algorithm 2.2 (TORSIONBASISFROMHINT).

MEDIUM

The C reference normalizes to $((A+2)/4 : 1)$ before basis generation. If the curve arrives from an isogeny codomain with $C_{24} \neq 1$, the Montgomery ladder produces a different projective representative, which propagates through DIFFERENCEPOINT and LADDER3PT.

Note in Algorithm 2.2 that the curve must be in normalized form before computing the basis.

8 Gluing requires Jacobian y -coordinates

Affects Algorithm 8.39 (GLUINGEVAL).

Low

Montgomery x -only arithmetic cannot distinguish $P + Q$ from $P - Q$. The gluing evaluation needs this distinction. The C reference lifts to Jacobian (x, y, z) via Okeya–Sakurai for this step—the only place in verification that needs y -coordinates. Note this in Algorithm 8.39.

9 Jacobian doubling for balanced strategy

Affects Algorithm 8.47, Phase 1.

Low

The spec does not specify whether Phase 1 doublings use Montgomery or Jacobian coordinates. The C reference doubles in Jacobian, then converts via `jac_to_xz`: $(x, y, z) \mapsto (x, z^2)$. The theta coordinate computations are sensitive to the projective representative. Note this, or at minimum state that the representative matters.

10 DIFFERENCEPOINT normalization factor

Affects Algorithm 8.12 (DIFFERENCEPOINT).

Low

The C reference multiplies B_{XX} , B_{XZ} , B_{ZZ} by $C \cdot \overline{C}^2 \cdot \overline{Z_P}^2 \cdot \overline{Z_Q}^2$ (Galois conjugation) before the quadratic solve. This makes the discriminant a fourth power in $\mathbb{F}_p$, reducing the $\mathbb{F}_{p^2}$ sqrt to an $\mathbb{F}_p$ sqrt. We initially used $(Z_P Z_Q)^2$, which lacks this property since $\overline{z}^2 \neq z^2$ in general. State the normalization factor and its purpose in Algorithm 8.12.

References

- [1] SQIsign team, *SQIsign: Compact Post-Quantum Signatures from Quaternions and Isogenies*, Specification version 2.0.1, dated 2025-07-07, <https://sqisign.org/spec/sqisign-20250707.pdf>.
- [2] M. N. H. Aardal, M. P. Azimova, E. L. Carrión, L. Dillinger, and M. J. Kannwischer, *Optimized One-Dimensional SQIsign Verification on Intel and Cortex-M4*, Cryptology ePrint Archive, Paper 2024/1563, <https://eprint.iacr.org/2024/1563>.
- [3] C. Costello, H. Hisil, and B. Smith, *Faster Compact Diffie–Hellman: Endomorphisms on the x -line*, ASIACRYPT 2017, <https://eprint.iacr.org/2017/518>.

A Never recompute $P-Q$ via DIFFERENCEPOINT

Affects Algorithm 4.9 (verification), lines 11–26, and transitively the $(2, 2)$ -isogeny chain. **CRITICAL**

A torsion basis $(P, Q, P-Q)$ is produced by TORSIONBASISFROMHINT (Algorithm 2.2). The third component $P-Q$ is computed once, via DIFFERENCEPOINT, which internally takes an $\mathbb{F}_{p^2}$ square root. Subsequent operations—scaling (repeated doubling), matrix application (Algorithm 4.9 line 14), and the even response isogeny (line 17)—must propagate all three points together. If at any point $P-Q$ is recomputed from P and Q via DIFFERENCEPOINT, the fresh square root may pick the opposite branch, producing a point that is *affinely distinct* from the original $P-Q$.

Concretely, the two branches of DIFFERENCEPOINT yield $x(P-Q)$ and $x(P+Q) = x(2P-Q)$; these are different points. All subsequent differential additions (the three-point ladder, the biladder for matrix application, and the $(2, 2)$ -isogeny chain’s LADDER3PT and LIFT_BASIS) consume $P-Q$ as a third input and produce wrong results if it is silently replaced by $2P-Q$.

In our implementation, three separate recomputations caused failures:

1. After scaling the basis by repeated doubling, we reconstructed $P-Q$ from the doubled P and Q instead of doubling $P-Q$ alongside them.
2. The matrix application $M_{\text{chl}} \cdot (P, Q)$ computed $P' - Q'$ via DIFFERENCEPOINT(P', Q') instead of a third biladder call $[(a-b)]P + [(c-d)]Q$.
3. Before entering the $(2, 2)$ -chain, we recomputed $P-Q$ from the post-isogeny images instead of pushing the original $P-Q$ through the isogeny as a third evaluation point.

Each of these independently caused KAT verification failure. Fixing all three was necessary and sufficient to make KAT vector 0 pass.

Current spec text. The spec does not warn against recomputing $P-Q$. Algorithms that scale or transform a basis mention only P and Q ; the third point $P-Q$ is not discussed.

Recommendation. Add a normative note in §2.2 or §8.2:

Once a basis $(P, Q, P-Q)$ has been produced by TORSIONBASISFROMHINT, the third point $P-Q$ must be propagated through every subsequent operation (doubling, matrix application, isogeny evaluation) alongside P and Q . Recomputing $P-Q$ via DIFFERENCEPOINT is not permitted: the square root may select a different branch, yielding an affinely distinct point that silently corrupts all downstream differential additions.

B Change-of-basis matrix application convention

Affects Algorithm 4.8 (SETCHANGEOFBASISMATRIX), Algorithm 4.9 (verification, line 14). **HIGH**

The spec does not specify whether the change-of-basis matrix $\mathbf{M}_{\text{chl}}$ is applied to a torsion basis (P, Q) by rows or by columns.

What the spec says. Algorithm 4.8 line 6 writes $(P_2, Q_2) \leftarrow \mathbf{M}_1 \cdot (P_2, Q_2)$ without specifying the convention. Algorithm 4.9 line 14 writes $(P_{\text{chl}}, Q_{\text{chl}}) \leftarrow \mathbf{M}_{\text{chl}} \cdot (P_{\text{chl}}, Q_{\text{chl}})$, again without specifying.

What the C reference does. `matrix_scalar_application_even_basis` in `verify.c` applies the matrix by *columns*:

$$\begin{aligned} R' &= [a] P + [c] Q & (\text{column 0: } \text{mat}[0][0], \text{mat}[1][0]) \\ S' &= [b] P + [d] Q & (\text{column 1: } \text{mat}[0][1], \text{mat}[1][1]) \end{aligned}$$

where $\mathbf{M} = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$. An implementor who applies by rows (the natural reading of “ $\mathbf{M} \cdot (P, Q)^T$ ”) will produce wrong results.

Impact. Applying by rows instead of columns silently produces the wrong torsion basis, causing the (2,2)-isogeny chain in verification to fail (typically a panic in the splitting step).

Recommendation. State explicitly in Algorithm 4.8 and Algorithm 4.9 that the matrix is applied by columns, or equivalently define the multiplication as $(P', Q') = (P, Q) \cdot \mathbf{M}^T$.

C GENERALIZEDREPRESENTINTEGER sampling ranges and divisibility

Affects Algorithm 3.12 (`GENERALIZEDREPRESENTINTEGER`), and transitively Algorithm 3.15 (`FIXEDDEGREEISOGENY`) during signing. MEDIUM

Algorithm 3.12 describes sampling z from $[1, \dots, m]$ (line 4) and t from $[-m', \dots, m']$ (line 5). Two aspects of the algorithm are easy to implement incorrectly:

1. Sign of t matters for the isogeny condition. Since $M' = 4M - p(z^2 + qt^2)$ depends on t^2 , it is tempting to sample only $t \geq 0$ and rely on the symmetry. However, the isogeny condition (line 15) checks $x - t \equiv 2 \pmod{4}$, which *does* depend on the sign of t . Restricting to $t \geq 0$ halves the effective search space for the isogeny condition, which can cause the algorithm to exhaust its bound without finding a solution.

2. Divisibility by 2 in O (line 18) is not coordinate-wise. Line 18 says “set d to be the biggest scalar such that $\gamma/d \in O$ ” and line 19 checks $d = 2$. For the standard order $O_0 = \mathbb{Z}[i] + j\mathbb{Z}[i]$ with $q = 1$ and $\omega = i$, an element $\gamma = x + iy + jz + kt$ satisfies $\gamma/2 \in O_0$ if and only if x, y, z, t are all even. However, for orders with half-integer basis elements (e.g., $O_0 = \mathbb{Z} + i\mathbb{Z} + \frac{1+j}{2}\mathbb{Z} + \frac{i+k}{2}\mathbb{Z}$, which is the actual O_0 for `SQIsign`), divisibility by 2 in O_0 is *not* equivalent to all coordinates being even in the $\{1, i, j, k\}$ basis. The correct check must account for the order’s common denominator and basis structure.

An implementor who reads “ $\gamma/d \in O$ ” as “all coordinates divisible by d ” (which is correct for $\mathbb{Z}^4$ but not for a non-trivial order) will write code that rejects valid solutions.

Recommendation. Add a remark after line 18 of Algorithm 3.12 clarifying:

The divisibility check “ $\gamma/d \in O$ ” is with respect to the order O , not coordinate-wise in the $\{1, i, j, k\}$ basis. For orders with a non-trivial denominator d_O , the integral coordinates of γ in the order’s column basis must all be divisible by $d \cdot d_O$.

Also note explicitly that the range $[-m', m']$ on line 5 includes negative values, and that both signs must be tried to satisfy the isogeny condition on line 15.

3. Product order in γ construction (line 17). Line 17 writes $\gamma \leftarrow (x + \omega y + jz + \omega jt)$. In quaternion algebra, $\omega j \neq j\omega$ in general ($ij = k$ but $ji = -k$). The C reference’s `quat_order_elem_create` computes $j \cdot \omega \cdot t$ (i.e., left-multiplying by j then by ω), which for $\omega = i$ gives $ji \cdot t = -kt$. Reading line 17 as $\omega jt = ij \cdot t = kt$ produces the opposite sign on the k component.

The difference does not affect the norm (`nrd` depends only on $|\text{components}|^2$), but it *does* affect the isogeny condition (line 15) and the divisibility check (line 18), since both depend on the parity and congruence class of the coordinates.

State whether line 17 means $(\omega j)t$ or $(j\omega)t$, or equivalently state the product order used by `quat_order_elem_create`.

D Response lattice: product vs. intersection

Affects Algorithm 4.2 (SQISIGN.SIGN), line 14.

High

Algorithm 4.2 line 14 writes

$$\alpha_{\text{rsp}} \leftarrow \text{RANDEQUIVALENTQUATERNION}(\overline{I_{\text{com}}} \cap I_{\text{sk}} \cdot I_{\text{chl}})$$

The notation “ $I_{\text{sk}} \cdot I_{\text{chl}}$ ” is ambiguous: it could mean the *ideal product* (the $\mathbb{Z}$ -span of all $\alpha\beta$ with $\alpha \in I_{\text{sk}}, \beta \in I_{\text{chl}}$) or the *intersection* $I_{\text{sk}} \cap I_{\text{chl}}$.

The C reference (sign.c:81–85) computes:

1. $I_{\text{chl,secret}} = I'_{\text{chl}} \cap I_{\text{sk}}$ (intersection, **not** product)
2. $\overline{I_{\text{com}}}$ (conjugate of I_{com})
3. $\alpha_{\text{rsp}} \leftarrow \text{SAMPLE}(I_{\text{chl,secret}} \cap \overline{I_{\text{com}}})$

Furthermore, the C ref intersects the *raw* challenge ideal I'_{chl} (before pushforward) with I_{sk} , not the pushed-forward $I_{\text{chl}} = [I_{\text{sk}}]_* I'_{\text{chl}}$.

An implementor who reads “ $I_{\text{sk}} \cdot I_{\text{chl}}$ ” as an ideal product will compute a completely different lattice. The ideal product’s norm is $\text{nrd}(I_{\text{sk}}) \cdot \text{nrd}(I_{\text{chl}}) \approx 2^{381}$, while the intersection’s norm divides $\text{lcm}(\text{nrd}(I_{\text{sk}}), \text{nrd}(I_{\text{chl}}))$. Sampling from the wrong lattice produces elements that generate trivial ideals (the HNF collapses to a scaled copy of $\mathcal{O}_0$), causing the response phase to fail silently.

Recommendation. Replace “ $I_{\text{sk}} \cdot I_{\text{chl}}$ ” with explicit “ $I_{\text{sk}} \cap I'_{\text{chl}}$ ” and note that the intersection uses the pre-pushforward challenge ideal.

E FIXEDDEGREEISOGENY missing u^{-1} normalization

Affects Algorithm 3.15 (FIXEDDEGREEISOGENY), and transitively Algorithm 3.13 (IDEALTOISOGENY) during signing.

CRITICAL

Algorithm 3.15 line 2 computes $\theta \leftarrow \text{REPRESENTINTEGER}(u \cdot (2^e - u), \mathcal{O}_t, \text{true})$, and line 4 constructs the $(2, 2)$ -isogeny kernel from $([u]P, \theta(P))$ and $([u]Q, \theta(Q))$. Since $\text{nrd}(\theta) = u \cdot (2^e - u)$, the endomorphism θ has degree $u(2^e - u)$; the kernel we want is that of the degree- $(2^e - u)$ isogeny, generated by $(P, (\theta/u)(P))$.

The C reference multiplies all four quaternion coordinates of θ by $u^{-1} \bmod 2^{e+2}$ before applying the endomorphism to the torsion basis (dim2id2iso.c, lines 133–143). Without this normalization, the ACTIONBYTRANSLATION determinant of the gluing kernel is degenerate: the splitting step finds zero (or ten) vanishing $U_{i,j}(0)$ indices instead of exactly one, and the $(2, 2)$ -chain fails.

Current spec text. Algorithm 3.15 line 4 writes “ $K_1 = (u \cdot P_t, \theta(P_t))$ ” without mentioning the u^{-1} normalization. The pseudocode is algebraically suggestive (the u -torsion is killed after $f-2-e$ doublings), but in practice the non-normalized kernel is numerically degenerate and no $(2, 2)$ -chain succeeds.

Recommendation. Replace line 4 of Algorithm 3.15 with $(P_t, (\theta \cdot u^{-1} \bmod 2^{e+2})(P_t))$ and note that u is odd (hence invertible $\bmod 2^{e+2}$) and that the extra $+2$ in the modulus accounts for the gluing’s order-4 structure.

F Fixed-precision Hermite Normal Form

Affects Algorithm 3.2 (HNF), and transitively every algorithm that constructs or reduces an ideal lattice — including Algorithm 3.9 (RANDEQUIVALENTPRIMEIDEAL), Algorithm 3.10 (RANDOMIDEALGIVENORM), and Algorithm 3.15 (FIXEDDEGREEISOGENY).

High

Algorithm 3.2 describes the column-style Hermite Normal Form via extended-GCD elimination. The pseudocode operates over $\mathbb{Z}$, i.e., arbitrary-precision integers.

What goes wrong. Classical HNF’s intermediate entries can grow exponentially in the rank. For rank-4 quaternion ideal lattices at NIST-I the initial column entries can reach $\sim 2^{770}$ bits (the $p \cdot g_i$ products from Algorithm 3.10). After a few xgcd elimination steps, intermediate column entries exceed any reasonable fixed-width budget. At our storage width of 1920 bits (30×64), the resulting “HNF” silently wraps: we observed basis columns collapse to elements of the ambient order O_0 (reduced norms of 4 and $\sim 2^{251}$) rather than the intended ideal I . Every subsequent operation — `RANDEQUIVALENTPRIMEIDEAL`, `IDEALTOISOGENY`, the entire response phase — then operates on a lattice unrelated to I .

What the C reference does. The C reference uses GMP (`ibz_t`), so coefficient growth is absorbed by arbitrary-precision arithmetic. The classical algorithm works correctly there.

Fix. Use the standard Domich–Kannan–Trotter modular HNF (Cohen, *A Course in Computational Algebraic Number Theory*, §2.4.2): given a positive multiple D of the lattice determinant $\det(L)$, reduce every intermediate entry modulo D after every arithmetic update. This bounds entries by D instead of by the exponential classical bound, at the cost of computing at a wider intermediate width $W \geq 2 \cdot \lceil \log_2 D \rceil / 64$. The result is identical to the classical HNF in $\mathbb{Z}$.

For NIST-I commitment ideals, $D = 4d^4N^2p \approx 2^{1282}$ is a compile-time constant. For response-phase and post-reduce ideals the modulus varies and is computed per call.

Recommendation. Add a remark after Algorithm 3.2 noting that fixed-precision implementations must account for coefficient growth, and that the standard modular HNF technique (reducing modulo a known multiple of the lattice determinant) is sufficient to keep intermediates bounded.

G Endomorphism image $P-Q$ and the Okeya-Sakurai lift

Affects Algorithm 3.15 (`FIXEDDEGREEISOGENY`), and transitively `IDEALTOISOGENY` and `KEYGEN`. **HIGH**

Algorithm 3.15 line 4 says to construct the kernel from $(\theta(P), \theta(Q))$. The $(2, 2)$ -isogeny chain lifts these Montgomery x -only points to Jacobian (x, y, z) coordinates via the Okeya-Sakurai formula, which requires the projective x -only representative of $P-Q$ (not just its affine x -coordinate).

The Okeya-Sakurai formula recovers y_Q from $(x_P, y_P, X_Q, Z_Q, X_{P-Q}, Z_{P-Q})$. The specific projective representative $(X_{P-Q} : Z_{P-Q})$ determines the recovered y_Q . Two representatives that agree in the ratio X/Z but differ in the individual coordinates produce different y_Q values—one correct, one not.

The spec does not specify how $\theta(P-Q)$ is obtained. Two natural approaches:

1. Compute $\theta(P-Q)$ as a third biladder call with scalars $(m_{00}-m_{01}, m_{10}-m_{11})$ on the same basis triple $(P, Q, P-Q)$.
2. Compute `DIFFERENCEPOINT`($\theta(P), \theta(Q)$), which takes an $\mathbb{F}_{p^2}$ square root and is ambiguous ($P-Q$ vs. $P+Q$).

Approach (1) produces three outputs whose projective representatives are consistent (all derived from the same basis via differential addition), so the Okeya-Sakurai lift recovers the correct y . Approach (2) picks a deterministic but potentially wrong branch of the square root, causing the lift to recover the wrong y for roughly half of all inputs. In our implementation, approach (2) caused the $(2, 2)$ -chain to fail for every θ from `REPRESENTINTEGER` while succeeding for the i automorphism (a degenerate case whose square root happens to land on the correct branch).

Recommendation. Add a remark in Algorithm 3.15 or in the description of the $(2, 2)$ -isogeny chain that the difference $\theta(P) - \theta(Q)$ must be computed by applying θ to $P - Q$ directly (via the biladder with the same basis triple), *not* by recomputing DIFFERENCEPOINT on the outputs. The same constraint applies to every other site that constructs a basis triple for the Okeya-Sakurai lift.

H GETINDEXSPLITTING must handle malformed input

Affects Algorithm 8.42 (GETINDEXSPLITTING), and transitively verification (Algorithm 4.9). **MEDIUM**

Algorithm 8.42 searches the null point $U_{i,j}(0)$ for the unique index (i, j) where it vanishes. The spec states this index “exists and is unique” (Proposition 8.5.3), which is true for honestly generated signatures.

What goes wrong. For *malformed* signatures (corrupted E_aux , M_chl , or chl), the $(2, 2)$ -isogeny chain may produce a null point with zero or multiple vanishing entries. An implementation that asserts uniqueness (e.g., via `debug_assert!`) will panic on such inputs, providing a denial-of-service vector: any untrusted signature can crash the verifier in debug builds. In release builds, the assertion compiles out and the function returns a wrong index, leading to silent misbehavior rather than a clean rejection.

We discovered this via negative test vectors (§H): single-bit flips in E_aux , M_chl , and chl all triggered the assertion.

Recommendation. Note in Algorithm 8.42 that implementations must handle the case where no (or multiple) vanishing indices exist, returning an error rather than assuming uniqueness. A verification implementation should treat a missing or ambiguous splitting index as signature rejection, not an assertion failure.

Signing key encoding requires caching the ideal generator

Severity. Low (implementation note).

Affected algorithm. Key encoding (§4.6).

Current spec text. The signing key wire format encodes the generator α of the secret ideal $I_{sk} = \mathcal{O}_0 \langle \alpha, N \rangle$ as four signed 32-byte coordinates. The ideal is defined by α and the norm N ; α is part of the key.

What goes wrong. An implementation that converts the ideal to HNF for internal use (as we do for lattice operations) loses the original generator. Recovering α from HNF requires a brute-force search over small linear combinations of basis vectors. The spec does not mention this recovery problem.

Recommendation. Note in §4.6 that implementations should cache α at parse/keygen time rather than attempting to recover it from the lattice representation.

I IDEALTOISOGENY line 6: matrix formula is under-specified

Affects Algorithm 3.13 (IDEALTOISOGENY), line 6.

HIGH

Algorithm 3.13 line 6 writes

$$[P, Q]^T \leftarrow \frac{1}{\text{nrd}(I) \text{nrd}(J_t)} M_{\beta_1^{-1}} M_{\beta_2} [\varphi_v(P_t), \varphi_v(Q_t)]^T.$$

The β_1^{-1} factor is a quaternion inverse, not a matrix inverse. The identity derived in the surrounding text,

$$\beta_2 \circ \varphi_{J_t} \circ \widehat{\varphi_{J_s}} \circ \widehat{\beta_1} = [\text{nrd}(I) \text{nrd}(J_s) \text{nrd}(J_t)] \widehat{\varphi_{I_2}} \circ \varphi_{I_1},$$

uses the conjugate $\widehat{\beta_1}$, and $\beta_1^{-1} = \widehat{\beta_1} / \text{nrd}(\beta_1)$ only as quaternions. An implementor writing $M_{\beta_1^{-1}}$ as a *matrix inverse* gets a different matrix from the intended $M_{\widehat{\beta_1}} / \text{nrd}(\beta_1) = \text{adj}(M_{\beta_1}) / \text{nrd}(\beta_1) \pmod{2^f}$ in general.

What the C reference does. `dim2id2iso.c:881–1014` builds the quaternion $\theta = \beta_2 \cdot \widehat{\beta_1}$ directly and scales its coordinates by $\text{invmod}(d_1 \cdot \text{nrd}(J_t), 2^f)$ before decomposing on $\mathcal{O}_0$. This is $\beta_2 \cdot \widehat{\beta_1}$, whereas the spec's $\beta_1^{-1} \beta_2 = \widehat{\beta_1} \beta_2 / \text{nrd}(\beta_1)$ has the factors in the opposite order; the two quaternion products agree only after the subsequent $\text{nrd}(\beta_1)$ scaling, and quaternion multiplication is not commutative in general.

Impact. An implementor has two independent traps: (a) interpreting $M_{\beta_1^{-1}}$ as the matrix inverse rather than the action matrix of β_1^{-1} viewed as a quaternion; (b) writing the quaternion product in the opposite order from the derivation.

Recommendation. Write line 6 with the conjugate factored out and the scaling collected:

$$[P, Q]^T \leftarrow \frac{1}{\text{nrd}(I) \text{nrd}(J_t) \text{nrd}(\beta_1)} M_{\widehat{\beta_1}} M_{\beta_2} [\varphi_v(P_t), \varphi_v(Q_t)]^T,$$

and state which quaternion-product order the matrix composition uses (equivalently, specify whether $M_{\alpha\beta} = M_\alpha M_\beta$ or $M_\beta M_\alpha$).

J IDEALTOISOGENY: β lives in the reduced ideal, not the input I

Affects Algorithm 3.13, Algorithm 3.16 (SUITABLEIDEALS), and anywhere the scaling $1/\text{nrd}(I)$ appears downstream of short-vector enumeration. HIGH

Algorithm 3.16 (SUITABLEIDEALS) returns β_i obtained from the short-vector enumeration inside $J_i I$, and $d_i = \text{nrd}(\beta_i) / \text{nrd}(J_i I)$. Efficient implementations (and the C reference, `dim2id2iso.c:534–565`) first replace I by the *smallest equivalent ideal* $I' = I \cdot \widehat{\delta} / \text{nrd}(I)$, so that the short-vector search produces usable degrees. The returned β and d are relative to I' , not to I , so the downstream invariant is

$$\text{nrd}(\beta) = d \cdot \text{nrd}(I') \neq d \cdot \text{nrd}(I).$$

What the C reference does. `dim2id2iso.c:670–684` “sends β back” to the original ideal after the pair is selected (multiplying by δ), so that by the time Algorithm 3.13 line 6 applies the scaling $1/\text{nrd}(I)$, the identity $\text{nrd}(\beta_1) = d_1 \cdot \text{nrd}(I)$ holds in terms of the caller-supplied I .

What goes wrong. An implementor who keeps β in the reduced ideal I' (a natural data representation) must also use $\text{nrd}(I')$ in every downstream scaling formula. Using the spec's $\text{nrd}(I)$ literal gives $\det(M_{\beta_1}) \neq d_1 \cdot \text{nrd}(I) \pmod{2^f}$, silently producing a non-isotropic outer kernel and `zeros=0` in the final splitting step.

Recommendation. State in Algorithm 3.16 that the returned β must be expressed in the input ideal's lattice (so $\text{nrd}(\beta) = d \cdot \text{nrd}(I)$), or name the lattice that β lives in and refer to its norm throughout Algorithm 3.13.

K `SUITABLEIDEALS` line 14: $v_2(u)$ vs. $v_2(\gcd(u, v))$

Affects Algorithm 3.16 (`SUITABLEIDEALS`), line 14.

MEDIUM

Algorithm 3.16 line 14 writes “ $e \leftarrow v_2(u)$ ” and line 15 divides both u and v by 2^e . Using $v_2(u)$ is only equivalent to $v_2(\gcd(u, v))$ when $v_2(u) \leq v_2(v)$:

- If $v_2(u) > v_2(v)$, dividing v by $2^{v_2(u)}$ loses low bits of v that were carrying real value.
- If $v_2(u) < v_2(v)$, dividing by $2^{v_2(u)}$ leaves v still even, so the next algorithm's precondition $\gcd(u \cdot d_1, v \cdot d_2) = 1$ fails.

What the C reference does. `dim2id2iso.c:833` computes `ibz_two_adic(&gcd(u, v))`, i.e. $v_2(\gcd(u, v))$.

Recommendation. Write line 14 as “ $e \leftarrow v_2(\gcd(u, v))$ ”.

L Algorithm 8.47 implementation variants are not specified

Affects Algorithm 8.47 (`ISOGENY22CHAINWITHTORSION`) and its callers, primarily Algorithm 3.13 line 9.
Low

The spec describes a single uniform chain of e 8-torsion isogenies with a kernel of order 2^{e+2} . The C reference implementation in `theta_isogenies.c` splits this into two variants, keyed on a boolean `extra_torsion`:

- `extra_torsion = true`: as in the spec, an e -length chain of 8-torsion isogenies.
- `extra_torsion = false`: an $(e - 2)$ -length chain of 8-torsion isogenies followed by a 4-isogeny and a 2-isogeny step. This consumes a kernel of order 2^e rather than 2^{e+2} , and the final two steps use square roots.

Algorithm 3.13 line 9 calls the outer chain with `extra_torsion = false` (`dim2id2iso.c:1128`) whereas the inner Algorithm 3.15 calls it with `extra_torsion = true` (`dim2id2iso.c:181`). The two variants produce the same $(2^e, 2^e)$ -isogeny but require different kernel torsions and different implementation paths.

Impact. Implementations that only supply the e -length 8-torsion variant must pad the kernel to order 2^{e+2} , which is impossible when e is close to the target curve's available 2-power torsion f . In particular, for $e \in \{f - 1, f\}$ the padded kernel does not exist. The C reference handles this by falling back to the shorter chain with 4-iso/2-iso tails, but the spec gives no guidance.

Recommendation. Either describe both variants in Algorithm 8.47 with a note on when each is required, or specify the 4-iso/2-iso tail as part of the algorithm so that implementations can support the full range of $e \leq f$ without ad-hoc workarounds.